

## PC-Trees vs. PQ-Trees

Wen-Lian Hsu

Institute of Information Science, Academia Sinica, Taipei

[Hsu@iis.sinica.edu.tw](mailto:Hsu@iis.sinica.edu.tw), <http://www.iis.sinica.edu.tw/IASL/hsu/eindex.html>

### Abstract

A data structure called *PC*-tree is introduced as a generalization of *PQ*-trees. *PC*-trees were originally introduced in a planarity test of Shih and Hsu [7] where they represent partial embeddings of planar graphs. *PQ*-trees were invented by Booth and Lueker [1] to test the consecutive ones property in matrices. The original implementation of the *PQ*-tree algorithms by Booth and Lueker using nine templates in each bottom-up iteration is rather complicated. Also the complexity analysis is rather intricate. We give a very simple linear time *PC*-tree algorithm with the following advantages: (1) it does not use any template; (2) at each iteration, it does all necessary tree-modification operations in one batch and does not involve the node-by-node bottom-up matching; (3) it can be used naturally to test the circular ones property in matrices; (4) the induced *PQ*-tree algorithm can considerably simplify Booth and Lueker's modification of Lempel, Even and Cederbaum's planarity test.

### 1. Introduction

A data structure called *PC*-tree is considered here as a generalization of *PQ*-trees. *PC*-trees were originally introduced to represent partial embeddings of planar graphs in Shih and Hsu [7]. *PQ*-trees were used to test the consecutive ones property in matrices [1]. However, the implementation of *PQ*-tree algorithms by Booth and Lueker [1] (hereafter, referred to as the *B&L algorithm*) is rather complicated. Also the complexity analysis is rather intricate. We shall present a very simple linear time *PC*-tree algorithm without using any template. Furthermore, we shall illustrate how to test the circular ones property in matrices using *PC*-trees.

*PQ*-trees were invented for the more general purpose of representing all permutations of a set  $U$  that are consistent with constraints of consecutiveness given by a collection  $C$  of subsets of  $U$  with the convention that the elements of each subset  $S$  in  $C$  must occur consecutively in the permutation.

The literature on problems related to *PQ*-trees is quite extensive. *PQ*-trees can be applied to test the consecutive ones property of (0,1)-matrices [1,3], to recognize interval graphs [1,2] and to recognize planar graphs efficiently [1,6]. Korte and Möhring [5] considered a modified *PQ*-tree and a simpler incremental update of the tree for the purpose of recognizing interval graphs. Klein and Reif [4] constructed efficient parallel algorithms for manipulating *PQ*-trees. On the other hand, *PC*-trees were initially used in Shih and Hsu [7] (*S&H*) to represent partial embeddings of planar graphs. It should be noted that *S&H*'s approach is entirely different from that of using *PQ*-trees in *B&L*'s modification [1] to Lempel, Even and Cederbaum's (*LEC*) planarity test [6]. *B&L* used *PQ*-trees to test the consecutive ones property of all nodes adjacent to the incoming node in their vertex addition algorithm. In *LEC*, the leaves of the *PQ*-tree must be those nodes adjacent to the incoming node. Internal nodes of the *PQ*-tree are not the original nodes of the

graph. Rather, they are there only to keep track of feasible permutations. Whereas in *S&H*, every *P*-node is an original node of the graph, every *C*-node represents a biconnected component in the partial embedding, and nodes adjacent to the incoming node can be scattered anywhere (both as internal nodes and as leaves) in the *PC*-tree. Thus, in *S&H*'s approach, a *PC*-tree faithfully represents a partial planar embedding of the given graph and is a more natural representation.

In this paper we shall focus on the application of *PQ*-trees to (0,1)-matrices. A (0,1)-matrix *M* has the *consecutive ones property (COP)* for the rows iff its columns can be permuted so that the ones in each row are consecutive. *M* is said to have the *circular ones property (CROP)* for the rows if its columns can be permuted so that either the ones or the zeros of each row are consecutive. When the zeros of a row are consecutive the ones will wrap around the first and the last column.

A *PQ*-tree is a rooted tree *T* with two types of internal nodes: *P* and *Q*, which will be represented by circles and rectangles, respectively. The leaves of *T* correspond 1-1 with the columns of *M*. Figure 1 gives an example of a matrix satisfying the *COP* and its corresponding *PQ*-tree.




Figure 1. A *PQ*-tree and the consecutive ones property

### 1.1. A Brief Introduction of the original *PQ*-Tree Algorithm of *B&L*

We shall follow the notations used in [1]. The frontier of a *PQ*-tree *T*, denoted by *FRONTIER*(*T*), is the permutation of the columns obtained by ordering the leaves of *T* from left to right. Such a permutation is called a *consistent permutation*. The *B&L* algorithm considers each row as a constraint and adds them one at a time. Each time a new row is added (a new iteration starts), the algorithm tries to modify the current *PQ*-tree to satisfy the consecutiveness constraint of columns of the newly added row. To guarantee unique *PQ*-tree representations we need to restrict ourselves to *proper PQ*-trees defined in [1]: every *P*-node has at least two children; every *Q*-node has at least three children. Two *PQ*-trees are *equivalent* iff one can be transformed into the other by applying zero or more equivalent transformations. The equivalence of two trees is written  $T \sim T'$ . There are two types of equivalent transformations:

1. Arbitrarily permute the children of a *P*-node,
2. Reverse the order of the children of a *Q*-node.

Denote the set of all consistent permutations (or frontiers) of a *PQ*-tree *T* by

$$\text{CONSISTENT}(T) = \{\text{FRONTIER}(T') \mid T' \sim T\}.$$

Given a *PQ*-tree *T* and a row *u* of *M*, define a new tree called the *u*-reduction of *T* to be one whose consistent permutations are exactly the original permutations in which the leaves in *u* occur consecutively. This new tree is denoted by *REDUCE*(*T*, *u*). Booth and Lueker gave a

procedure *REDUCE* to obtain the  $u$ -reduction. The procedure applies a sequence of *templates* to the nodes of  $T$ . Each template has a *pattern* and a *replacement*. If a node matches the template's pattern, the pattern is replaced within the tree by the template's replacement. This is a bottom-up strategy that examines the tree node-by-node obeying the child-before-parent discipline. In Section 3, we shall describe a more global strategy that does the tree-modification operation in one batch operation at each iteration.

A node is said to be *full* if all of its descendant leaves are in  $u$ ; It is *empty* if none of its descendant leaves are in  $u$ ; If some but not all of the descendant leaves are in  $u$ , it is said to be *partial*. Nodes are said to be *pertinent* if they are either full or partial. The pertinent subtree of  $T$  with respect to  $u$ , denoted by  $PERTINENT(T, u)$ , is the subtree of minimum height whose frontier contains all columns in  $u$ . The root of the pertinent subtree is denoted by  $ROOT(T, u)$ , which is usually not the root of the entire tree. There are nine templates used in [1] as shown in Figure 9. Two passes are used in the actual implementation. The first identifies the nodes to be processed and the second applies the templates.

A naïve implementation would classify each node as empty, full, or partial and then perform a pattern replacement. This will run into the danger of processing the entire tree even if the sets being reduced are small at each iteration. In order to achieve optimal efficiency, *B&L* algorithm takes several precautions in scanning the pertinent subtree. For example, the maintenance of parent pointers may cause a problem. It is possible that almost all nodes in the tree may receive a new parent even though the row  $u$  has only two elements. Therefore, parent pointers are only kept for children of  $P$ -nodes and for endmost children of  $Q$ -nodes. Interior children of a  $Q$ -node do not keep parent pointers. Rather, they borrow the pointers from their endmost siblings. In *B&L* algorithm, they adopted the notion of blocked, unblocked nodes and a block of blocked nodes in dealing with the borrowing of parent pointers. These measures unavoidably complicate the implementation of the algorithm. In Section 5, we shall adopt an alternative strategy of “parallel search”, which identifies a parent pointer for each node immediately and costs at most  $3lu$  collectively at each iteration. Parallel search considerably simplifies the implementation.

## 1.2. Definition of *PC*-Trees

A *PC*-tree is a rooted tree  $T$  with two types of internal nodes:  $P$  and  $C$ , which will be represented by circles and double circles, respectively. The leaves of  $T$  correspond 1-1 with the columns of  $M$ . Figure 2 gives an example of a matrix satisfying the *CROP* (but not the *COP*) and its corresponding *PC*-tree. The frontier of a *PC*-tree  $T$ , denoted by  $FRONTIER(T)$ , is the circular permutation of the columns obtained by ordering the leaves of  $T$  in a counter-clockwise circular order. Such a permutation is called a *consistent circular permutation*. Similar to the *PQ*-tree notations, two *PC*-trees are said to be *equivalent* iff one can be transformed into the other by applying zero or more equivalent transformations. The equivalence of two trees is written  $T \sim T'$ . There are two types of equivalent transformations:

- 1、 Arbitrarily permute the children of a  $P$ -node,
- 2、 Change the circular order of the neighbors of a  $C$ -node from clockwise to counter-clockwise.

Denote the set of all consistent circular permutations (or frontiers) of a *PC*-tree  $T$  by  $CONSISTENT(T) = \{FRONTIER(T') \mid T' \sim T\}$ . Given a *PC*-tree  $T$  and a row  $u$  of  $M$ , define a new tree called the  $u$ -reduction of  $T$  to be one whose consistent circular permutations are exactly the original permutations in which the leaves in  $u$  occur consecutively in the circular list. This

new tree is denoted by  $REDUCE(T, u)$ .




Figure 2. A  $PC$ -tree and the circular ones property

For each  $PQ$ -tree, we can obtain its corresponding  $PC$ -tree by replacing each  $Q$ -node  $w$  with a  $C$ -node and merging the parent  $x$  of  $w$  into the left-right ordering of the children of  $w$  (make  $x$  a neighbor of the two endmost children of  $w$ ) to form a circular ordering. Note that flipping a  $Q$ -node in a  $PQ$ -tree is the same as changing the circular neighbor list of its corresponding  $C$ -node from clockwise to counter-clockwise order. There are a few differences between  $PC$ -trees and  $PQ$ -trees: (1) When a  $PQ$ -tree is transformed into a  $PC$ -tree, and its root is a  $Q$ -node, then the two endmost children of the root  $Q$ -node must remain endmost in any left-right ordering in order to maintain the  $COP$  in any consistent circular permutation of the transformed  $PC$ -tree. Aside from this restriction, all other  $PC$ -tree operations will yield equivalent  $PQ$ -tree operations; (2) Since it is the circular order of the leaves that needs to be preserved (for example, in Figure 2, we can rotate the tree so that  $FRONTIER$  becomes 2,3,4,5,6,1), one might as well fix the first column to be 1 in a  $PC$ -tree in considering any consistent circular permutation; (3) Unlike a  $PQ$ -tree, the root of a  $PC$ -tree in general is not essential except to maintain the child-parent relationships.

Our improvement over the  $B\&L$  algorithm is based on a simple observation that there exist at most two special partial nodes and the unique path connecting them gives a streamline view of the update operation. In particular, no template is necessary in our  $PC$ -tree algorithm and the node-by-node template examination is replaced by *one* swift batch operation. Compared to the two-pass strategy of  $B\&L$ , one might view our approach as a one-pass strategy that involves only node labeling operation. The second pass is replaced by a trivial tree-replacement operation.

In summary, our  $PC$ -tree algorithm has the following advantages: (1) it does not use any template; (2) at each iteration, it does all necessary tree-modification operations in one batch and does not involve the node-by-node bottom-up matching; (3) it can be used naturally to test the circular ones property in matrices; (4) the induced  $PQ$ -tree algorithm can considerably simplify  $B\&L$ 's modification of Lempel, Even and Cederbaum's planarity test.

## 2. A Forbidden Structure of Consistent Permutations

A key observation that simplifies our  $PC$ -tree algorithm is based on the following forbidden structure of consistent permutations.

**Theorem 2.1.** The following structure is forbidden in any consistent circular permutation of the columns of  $M$  that preserves the  $CROP$ . Let  $S_1 = \{a_1, b_1\}$ ,  $S_2 = \{a_2, b_2\}$  and  $S_3 = \{a_3, b_3\}$  be three subsets of columns such that

- (1) Columns  $a_1, a_2, a_3, b_1, b_2$  and  $b_3$  are distinct from each other.
- (2) No two columns in any  $S_i$  are separated by columns in the other two subsets in any consistent

circular permutation. Namely, columns cannot be arranged in the following order, where  $j, k$  are distinct from  $i$  (but  $j$  could be equal to  $k$ ):

(a column in  $S_i$ )...(a column in  $S_j$ )...(a column in  $S_i$ )...(a column in  $S_k$ )...

- (3) Furthermore,  $b_1, b_2$  and  $b_3$  are not separated by any of the  $a_i$ 's, namely, columns must be arranged in the following order, where  $i, j, k$ , (similarly,  $p, q, r$ ) are all distinct from each other.

$a_i \dots a_j \dots a_k \dots b_p \dots b_q \dots b_r \dots$

**Proof.** From (2), we must have the following arrangement  $a_1 \dots b_1 \dots b_2 \dots a_2 \dots$  (or the reverse) in any consistent circular permutation of the columns. By (2) and (3), one can only place  $a_3, b_3$  into the above arrangement as follows:

- I.  $a_1 \dots a_3 \dots b_3 \dots b_1 \dots b_2 \dots a_2 \dots$
- II.  $a_1 \dots b_1 \dots b_2 \dots b_3 \dots a_3 \dots a_2 \dots$

However, in (I), columns in  $S_1$  are separated by other subsets; and in (II), columns in  $S_2$  are separated by other subsets. ■

### 3. Terminal Nodes and Paths

Note that the definition of empty, partial and full nodes are all relative to a fixed root in the tree  $T$ . In the following we shall define a notion that is independent of the selection of the root of  $T$ . Pick any node  $w$  of  $T$  as a root, one can form a tree rooted at  $w$ , denoted by  $T(w)$ . Define a node  $t$  of  $T$  to be a *terminal node* if there exists a node  $r$  in  $T$  such that, in  $T(r)$ ,  $t$  is partial and none of  $t$ 's children is partial. In other words, each child of a terminal node is either empty or full relative to a particular rooted tree. As long as there is a partial node, there must exist a descendent terminal node. Terminal nodes play a major role in simplifying our tree-update operation.

**Theorem 3.1.** If the given matrix  $M$  satisfies the *CROP*, then there exist at most two terminal nodes in  $T$  at every iteration.

**Proof.** Suppose, on the contrary, there are three terminal nodes  $t_1, t_2$ , and  $t_3$ . Let  $r_1, r_2, r_3$  be the roots, respectively, that make them terminal nodes. We claim that none of the subtree  $T_{t_i}$  with root  $t_i$  within  $T(r_i)$  can contain the other two terminal nodes. It suffices to show that the subtree  $T_{t_1}$  with root  $t_1$  in  $T(r_1)$  does not contain  $t_2$ .

Suppose, on the contrary,  $T_{t_1}$  contains  $t_2$ . Let  $ST_{t_2}$  be the subtree with root  $t_2$  within the subtree  $T_{t_1}$ . Since  $t_2$  is either full or empty in  $T_{t_1}$ ,  $r_2$  must be contained in  $ST_{t_2}$  (for otherwise, the subtree at  $t_2$  within the tree rooted at  $r_2$  would still be empty or full). Without loss of generality, assume  $t_2$  is full in  $T_{t_1}$ .

Consider the parent  $w$  of  $t_2$  in  $T_{t_1}$ . If  $w = t_1$ , then we have a case that  $t_2$  is full in  $T(r_1)$  and  $t_1$  is empty in  $T(r_2)$  and  $t_1, t_2$  are adjacent. Then, we can conclude: (1) if  $r_3$  is contained within the subtree rooted at  $t_1$  in  $T(r_2)$ , then  $t_1$  would be the only terminal node; and (2) if  $r_3$  is contained within the subtree rooted at  $t_2$  in  $T(r_1)$ , then  $t_2$  would be the only terminal node. In both cases, we cannot have a distinct third terminal node  $t_3$ .

Hence, we can assume  $w \neq t_1$ . We claim that  $w$  would become a partial child of  $t_2$  in  $T(r_2)$  as shown in Figure 3 (due to the reversal of parent-child relationship). This would be contradictory to the fact that  $t_2$  is a terminal node in  $T(r_2)$ . Note that making the tree rooted at  $r_2$  would only reverse the parent-child relationships for nodes along the path from  $r_2$  to  $r_1$  in  $T(r_1)$ . To prove this claim, let  $x (\neq t_2)$  be another child of  $w$  in  $T_{t_1}$ . Since  $t_1$  is a terminal node in  $T_{t_1}$ ,  $w$  must also be full. Hence,  $x$  is full in  $T_{t_1}$ . Furthermore,  $x$  will remain full in  $T(r_2)$ . However,  $t_1$  must have an empty

child  $y$  in  $T_{t_1}$ , which will remain empty in  $T(r_2)$ . Since both  $x$  and  $y$  are descendants of  $w$  in  $T(r_2)$ ,  $w$  must be partial in  $T(r_2)$ .

Now, choose an empty child  $a_i$  and a full child  $b_i$  from each of the  $t_i$ ,  $i = 1, 2, 3$ . Then,  $S_1 = \{a_1, b_1\}$ ,  $S_2 = \{a_2, b_2\}$  and  $S_3 = \{a_3, b_3\}$  constitute a forbidden structure as described in Theorem 2.1.




Figure 3. The tree rooted at  $r_1$ , the subtrees  $T_{t_1}$  and  $ST_{t_1}$

In case there is only one terminal node  $t$ , then in the tree rooted at  $t$ , every child of  $t$  is either full or empty. If  $t$  is a  $C$ -node, then the full children of  $t$  must be consecutive in its circular list and there is no tree modification necessary. If  $t$  is a  $P$ -node, then create a new  $P$ -node  $w$  as a child of  $t$  and change the parent of all full children of  $t$  to  $w$ .

The more interesting case occurs when there are two terminal nodes  $t_1$  and  $t_2$ . Define the unique path in  $T$  connecting  $t_1$  and  $t_2$  to be the *terminal path*. Now, any  $C$ -node in the interior of the terminal path has its children (not on the path) divided into two sides. Since one can flip the children ordering (change from clockwise to counter-clockwise) along the terminal path to obtain an equivalent  $PC$ -tree, we shall show that there is a unique way to flip the children of these  $C$ -nodes “correctly”. Denote  $ROOT(T, u)$  by  $m$ .

**Theorem 3.2.** If the matrix  $M$  satisfies the *CROP* and there are two terminal nodes  $t_1$  and  $t_2$  at some iteration  $u$ , then

- (1) Every  $C$ -node in the interior of the terminal path satisfies that its full (respectively, empty) children together with its two neighbors on the path are contiguous in its circular list. Therefore, they can be flipped to one side correctly.
- (2) If a terminal node, say  $t_2$ , is a  $C$ -node, then consider two cases: 1.  $t_2$  is not on the unique path from the root to  $t_1$ . Then its full children together with its parent are contiguous in its circular list and exactly one of its parent’s siblings is full; 2.  $t_2$  is on that path. If  $t_2$  is the root, then its full children together with its unique child, say  $w$ , on the path are contiguous in its



Figure 5. The terminal path with empty nodes flipped away from the root

### 3 、 Constructing the New *PC*-tree

Rather than using the node-by-node tree modification in the *B&L* algorithm, we shall update the current *PC*-tree to the new one in a batch fashion. In the following we divide our algorithm into two stages: the labeling operation, and the tree-update operation. The latter consists of the splitting operation and the modifying operation.

When a new row  $u$  is added, perform the following labeling operation:

1. Label all leaf columns in  $u$  *full*.
2. The first time a node becomes partial or full, report this to its parent.
3. The first time a node  $x$  gets a partial or full child label  $x$  *partial*
4. The first time all children of a node  $x$  become full label  $x$  *full*

Our tree-update batch operation consists of the following steps:

1. Delete all edges on the terminal path
2. The *splitting operation*: duplicate the nodes on the terminal path to create a cyclic order as shown in Figure 6. Contract all degree two nodes
3. The *modifying operation*: create a new *C*-node  $w$  and connect  $w$  to duplicated *P*-nodes and *C*-nodes according to their relative positions on the terminal path as shown in Figure 7. The detail is as follows: connect  $w$  directly to all *P*-nodes that have nonempty leaves; connect  $w$  to all full children of *C*-nodes according to their original ordering as shown in Figure 8.

In fact, by considering steps 2 & 3 together, we can see that every *P*-node on the path gets duplicated and every *C*-node on the path is deleted.

The following figures illustrate the splitting and the modifying operations on the graph of figure 3(a). Note that the root in a *PC*-tree does not play a major role in these operations.




Figure 6. The splitting operation






Figure 7. Connecting to the new C-node






Figure 8. The modifying operation

**Theorem 5.** The corresponding  $PQ$ -tree of the newly constructed  $PC$ -tree is the same as the one produced by the  $B\&L$  algorithm.


**Proof.** As noted before,  $B\&L$  algorithm modifies the tree in a node-by-node bottom-up fashion based on the matching of 9 templates. Whenever a pattern is matched, it is substituted by the replacement. The templates used in  $B\&L$  algorithm are illustrated in Figure 9. We shall skip the easy cases of  $P0$  and  $P1$ ,  $Q0$  and  $Q1$  [1].




Template  $P2$  for  $ROOT(T,S)$  when it is a  $P$ -node



Template  $P3$  for a singly partial  $P$ -node which is not  $ROOT(T,S)$



Template  $P4$  for  $ROOT(T,S)$  when it is a  $P$ -node with one partial child ( $T,S$ )



Template  $P5$  for a singly partial  $P$ -node, other than  $ROOT(T,S)$ , with one partial child

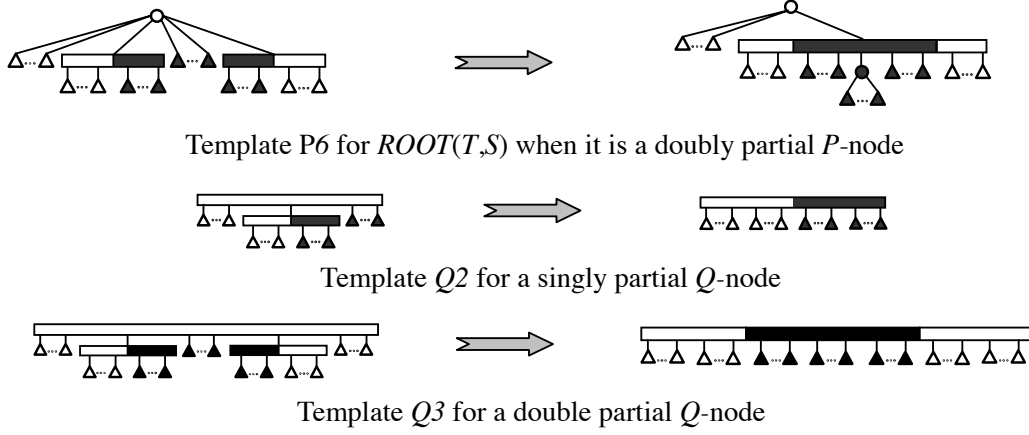


Figure 9. The template operations of  $B\&L$  algorithm

We shall prove the theorem by induction on the depth of the pertinent subtree at which  $B\&L$  executed the template operation. For each template operation of the  $PQ$ -tree, we shall illustrate the corresponding splitting operation of the  $PC$ -tree and show the equivalence of the two operations in terms of the resulting  $PQ$ -tree. The basis of the induction at depth 1 occurs when the operation is performed on the children of  $ROOT(T,u)$ .

In Figures 10 to 14 below we use the following convention: The two leftmost diagrams give a  $PQ$ -tree (on the top) and its corresponding  $PC$ -trees (at the bottom). The top part of each figure shows the  $PQ$ -tree replacement. The bottom part gives the corresponding  $PC$ -tree splitting and modifying operation. The equivalence is shown by the two rightmost diagrams, in which a  $PQ$ -tree (on top) and its corresponding  $PC$ -tree (in the bottom) are obtained through  $B\&L$  template matching and the  $PC$ -tree operation, respectively.

We first consider templates applied at  $ROOT(T,u)$ . In Figure 10 we consider template P2.




Figure 10. Template P2 for  $ROOT(T,S)$  when it is a  $P$ -node

The same goes for the templates P4, P6, Q2 and Q3 as shown in Figures 9,10,11,12.




Figure 11. Template  $P4$  for  $ROOT(T,S)$  when it is a  $P$ -node with one partial child  $(T,S)$




Figure 12. Template  $P6$  for  $ROOT(T,S)$  when it is a doubly partial  $P$ -node

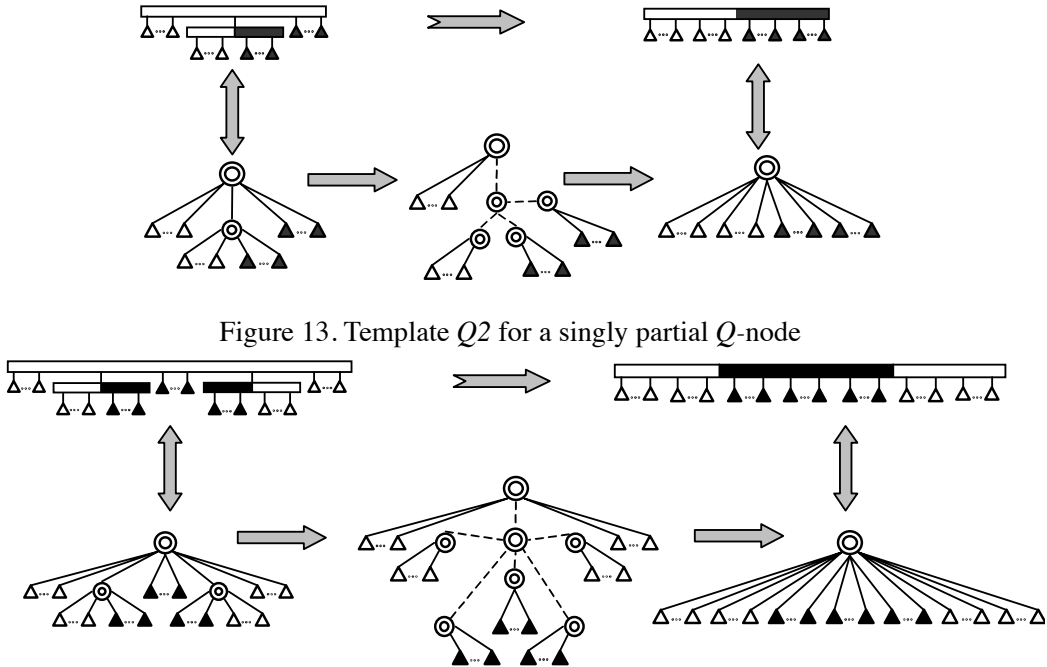


Figure 13. Template  $Q2$  for a singly partial  $Q$ -node

Figure 14. Template  $Q3$  for a double partial  $Q$ -node

The above has established the basis for the induction. Now, assume the theorem is true for

$PQ$ -trees whose pertinent subtree is of depth  $k$  and consider a tree whose pertinent subtree is of depth  $k+1$ . In the  $B\&L$   $PQ$ -tree operation, a template will be applied to a non-root terminal node  $x$  of the current  $PQ$ -tree  $T_1$ , to obtain a new  $PQ$ -tree  $T_2$ . To simplify the argument we assume the pertinent subtree of  $T_2$  has depth  $k$ . Since further  $PQ$ -tree operations on  $T_2$  can be done equivalently through the corresponding  $PC$ -tree operations by induction, we only have to show that the  $PC$ -tree operation performed on  $T_1$  results in the same  $PC$ -tree performed on  $T_2$ . In Figures 15 and 16 below we describe the  $PC$ -tree operations: (a) Transform the current  $PQ$ -tree into a  $PC$ -tree; (b) Illustrate the  $PC$ -tree operation on the entire tree. However, the part that is related to this particular node is highlighted with the remaining part (which is the same for all diagrams) indicated by thick dotted line. Note that, with the bottom-up strategy of  $B\&L$  algorithm, the node in consideration must be a terminal node in the  $PC$ -tree, which is at the very end of the splitting operation. (c) The final modifying operation where the newly created  $C$ -node, say  $w$ , is connected directly to all  $P$ -nodes that have nonempty leaves and  $w$  is connected to all full children of  $C$ -nodes according to their original ordering (the original  $C$ -nodes on the terminal path are deleted). The original  $PQ$ -tree templates are shown by the leftmost two diagrams.




Figure 15. Template  $P3$  for a singly partial  $P$ -node which is not  $ROOT(T,S)$

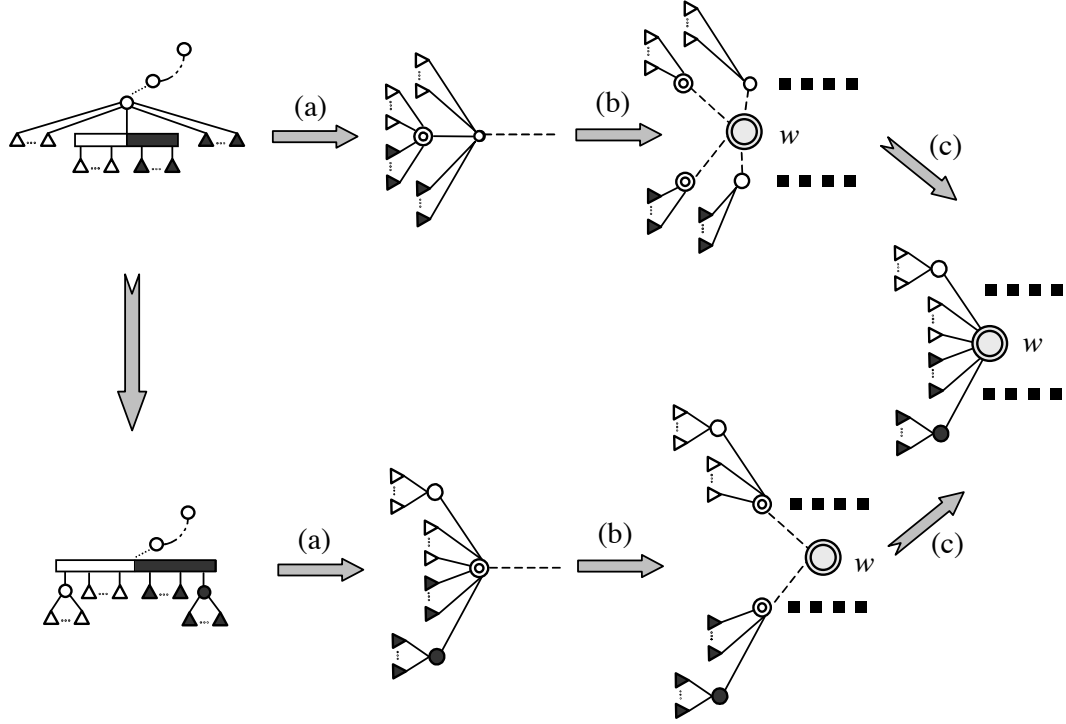


Figure 16. Template *P5* for a singly partial *P*-node, other than  $ROOT(T, S)$ , with one partial child  
**End of Proof of Theorem 4.1.** ■

**Corollary 4.2.** The *PC*-tree algorithm can be used to test the circular ones property.

**Proof.** Theorem 4.1 compared the *PC*-tree algorithm versus the *PQ*-tree algorithm concerning the consecutive ones testing and showed that they result in the same *PQ*-tree at each iteration. For the *CROP* test, the difference is that, at some iteration, we might encounter a row  $u$  whose complementary column set satisfies the *COP*. By Theorem 3.1, one can still obtain the terminal nodes and the terminal path for such a row. However, at this time, the empty nodes (rather than the full nodes) are flipped away from the root as illustrated in Figure 4. Since our splitting and modification operations are completely independent of the root of the tree, such a change does not affect our *PC*-tree modification.

On the other hand, *B&L* algorithm uses templates based on the consideration of pertinent nodes. Therefore, they could not afford to consider templates based on the consideration of empty nodes (which is exactly what is needed in this case if one adopts the template matching strategy). Hence, their implementation cannot be used to test the *CROP* directly and a transformation on the matrix itself is needed to achieve that [3]. ■

## 5. Implementation Details

In the last paragraph of Section 1.1, we mentioned an important implementation trick in *B&L*'s *PQ*-tree algorithm -- the borrowing of parent pointers. We shall do the same in our *PC*-tree implementation. Namely, parent pointers are only kept for children of *P*-nodes and for endmost children of *C*-nodes; parent pointers for internal children must be borrowed from the endmost ones through their siblings. However, in identifying the parent for children of a *C*-node, we shall

adopt the following “parallel search” strategy, which considerably simplifies the implementation.

The notion of parallel search is defined within a circular linked list of a  $C$ -node. Let  $C_1$  be a  $C$ -node and  $t$  be a node (which does not have a parent pointer) in the circular linked list of  $C_1$ . A *parallel search (PS)* from  $t$  is a traversal in the circular linked list from  $t$  along both the clockwise and the counter-clockwise directions in equal pace. Every node traversed will be marked. The purpose of the *PS* is to identify the parent for  $t$ , namely  $C_1$ . Let  $v_1$  and  $v_2$  be the two endmost children of  $C_1$  that have parent pointers to  $C_1$ . The very first *PS* in  $C_1$  from a node will terminate when one of the endmost children is reached. Afterwards, a *PS* will terminate when an endmost child is reached or a previously searched node is reached. In both cases, the parent for  $t$  can be identified. Since the time spent in each iteration should be proportional to the number of incoming columns,  $lul$ , we shall terminate a *PS* after  $lul$  steps.

There are three cases that could occur at the end of the labeling operation for the node  $C_1$ : (1) all of its children are labeled; (2) a contiguous subset of its children including one of its endmost children are labeled; (3) a contiguous subset of its children are labeled but none of its endmost children is labeled. In both cases (1) and (2), the total cost of all *PS*s occurred for the node  $C_1$  is proportional to the number of labeled children of  $C_1$ . Case (3) can occur only when  $C_1$  is  $ROOT(T, u)$ , which incur a cost of  $lul$  and can happen at most once in the iteration for  $u$ . Therefore, if the cost for parallel search is charged separately, we could assume each child of  $C_1$  can get a parent pointer in constant time. This is exactly what we have assumed in the labeling operation in Section 4.

Now, consider the tree-update operation.  $C$ -nodes on the terminal path are deleted and  $P$ -nodes are duplicated. Each node on the path has at least one child that contains a distinct leaf column in  $lul$ . Hence, the number of duplicated nodes (or the number of nodes on the terminal path plus one - the newly created  $C$ -node) is no greater than  $lul + 1$ . The sibling relationships of the new  $C$ -node can be constructed by connecting  $P$ -nodes and the neighboring children of  $C$ -nodes on the path according to their order on the path (as shown in Figure 17), which takes time proportional to the number of nodes on the path (and bounded by  $lul$ ).




Figure 17. The construction of the new sibling relationships

## 6. Complexity Analysis

Since the number of new nodes generated is  $O(lul)$ , the number of steps involved in both the labeling operation and the tree-update operation are bounded by  $O(lul)$ , the running time at iteration  $u$  is bounded by  $O(lul)$ . Hence, the time complexity of the algorithm is linear.

Compared to the *B&L* algorithm, our major saving is two folds: first, we have avoided the pattern-matching phase at each node; secondly, our parallel search strategy has considerably simplified the borrowing of parent pointers for internal children of  $C$ -nodes.

## 7. Acknowledgement

We are grateful for the generous support from the National Science Council under Grant NSC 90-2213-E-001.

### References

1. K.S. Booth and G.S. Lueker, *Testing of the Consecutive Ones Property, Interval graphs, and Graph Planarity Using PQ-Tree Algorithms*, **J. Comput. Syst. Sci.** 13, 3 (1976), 335-379.
2. D.R. Fulkerson and O.A. Gross, *Incidence Matrices and Interval Graphs*, **Pacific Journal of Math.**, (1965), 15:835-855.
3. M.C. Golumbic, **Algorithmic Graph Theory and Perfect Graphs**, Academic Press, New York, 1980.
4. P.N. Klein and J.H. Reif, *An efficient parallel algorithm for planarity*, **J. of Computer and System Science** 37, (1988), 190-246.
5. N. Korte and R. H. Möhring, *An Incremental Linear-Time Algorithm for Recognizing Interval Graphs*, **SIAM J. Comput.** 18, 1989, 68-81.
6. A. Lempel, S. Even and I. Cederbaum, *An Algorithm for Planarity Testing of Graphs*, **Theory of Graphs**, ed., P. Rosenstiehl, Gordon and Breach, New York, (1967), 215-232.
7. W.K. Shih and W.L. Hsu, *Note A new planarity test*, **Theoretical Computer Science** 223, (1999), 179-191.